

25 Random number generators

(25.1) What is a random number? **In theory**, a random number, or a random sequence of numbers, is one which cannot be produced by a program shorter than the number itself (Kolmogorov complexity).

In practice, we use sequences of numbers which are produced by extremely short programs! By the Kolmogorov criterion, they are certainly *not* random. These should be called *pseudo-random*, but the ‘pseudo-’ is generally forgotten.

(25.2) Until recent years, the following method, linear congruential random number generator, was used.

$$X_{n+1} = (aX_n + c) \bmod m$$

Here, a, c, X_0 , and m are constants, chosen suitably. Typically $m = 2^{31}$. In `stdlib.h`, `RAND_MAX` (which is $m - 1$) is given as $2^{31} - 1$, over 2 billion.

`man rand`

it will tell you a little more (probably inaccurate, judging by a look at `stdlib.h`). The numbers X_0, X_1, X_2, \dots are returned by successive calls to `rand()` (which remembers the previous number in a static variable).

(25.3) **The sequence is always the same.** The main point is that it ‘looks’ random under various statistical and empirical tests.

(25.4) Obviously some choices of constant are bad. For example, if $a = 0$ then $X_n = c$ for all n . If $a = 1$ then $X_n = (X_0 + nc) \bmod m$, which neither looks random nor is. We may assume that $a \geq 2$.

In this case it is easily shown by induction on n that

$$X_n = (a^n X_0 + \frac{a^n - 1}{a - 1} c) \bmod m.$$

(25.5) Since the numbers are between 0 and $m - 1$, eventually we get a repeated value, so for some $n \geq 0$ and $d > 0$, $X_n = X_{n+d}$. The smallest such d is called the *period*.

(25.6) **Using genuinely random seeds.** A simple method is to read the system clock, getting the microsecond part of the time. The chances against getting the same seed twice are about a million to one. **Don’t reset the seed** in the middle of a program. Refer to the programming notes.

(25.7) **Producing various random distributions.** Use `man drand48`. Most often, we want a random number between 0 and $b - 1$ for some smallish bound b . The obvious method (which I have used in ignorance for many years) is `rand() % b`. The manuals warn us not to do this. The reason is simple. Suppose a linear congruential generator is used. Suppose b divides m . Then one is effectively using the generator

$$X_{n+1} = (aX_n + c) \bmod b$$

whose period is $\leq b$: hardly a good choice.

The recommended method is, for example, to take the double-precision number `b * drand48()` and round it down to the nearest integer. If `d` is an `int`, then

```
d = b * drand48();  
d = d % b;
```

will accomplish this. This belongs to the programming notes.

(25.8) Linear congruential generators have the property that if m is a power of 2 (which it usually is), then the low-order bits are not ‘random’. (This follows from the previous paragraph.)

25.1 Example: estimating π

```
#include <stdio.h>  
#include <stdlib.h>  
#include <sys/time.h>  
  
double randval ()  
{  
    static int first = 1;  
    if ( first )  
    {  
        struct timeval tv;  
        gettimeofday ( & tv, NULL );  
        srand48 ( tv.tv_usec );  
        first = 0;  
    }  
  
    return drand48 ();  
}  
  
main( int argc, char * argv [] )  
{  
    int n = atoi ( argv[1] );  
  
    int i;  
    double sum = 0;  
    for (i=0; i<n; ++i)  
    {  
        double x = randval();  
        double y = randval();  
        if ( x*x + y*y <= 1 )  
            sum += 1;  
    }  
    printf("%d trials, pi is %f\n", n, 4*sum/n );  
}
```

```

% a.out 10000
10000 trials, pi is 3.117600
% a.out 10000
10000 trials, pi is 3.147200
% a.out 10000
10000 trials, pi is 3.160400

```

There is a special-purpose technique for generating *normally distributed* random numbers. Of course, the usual random number generator produces random numbers in the range $[0, 1)$ (so 1 never happens). Actually, 0 would break the program.

```

#include <stdio.h>
#include <ctype.h>
#include <math.h>
#include <stdlib.h>
#include <sys/time.h>

// Marsaglia method
// U,V uniform on unit disc
// scale by sqrt((-ln S)/S) where S=U^2+V^2

void two_normally( double * x, double * y )
{
    double v1, v2, s, mul;

    int found = 0;

    while ( ! found )
    {
        v1 = 2 * drand48() - 1;
        v2 = 2 * drand48() - 1;
        s = v1*v1 + v2*v2;
        found = ( s < 1 );
    }

    mul = sqrt ( - 2 * log(s) / s );
    *x = v1 * mul;
    *y = v2 * mul;
}

main ( int argc, char * argv[] )
{

```

```

struct timeval tv;
int i,n;

n = atof ( argv[1] );
double sum = 0, sumsquares = 0;

gettimeofday ( &tv, NULL );
srand48 ( tv.tv_usec );

double x,y;
for (i=0; i<n; i += 2)
{
    two_normally ( &x, &y );
    sum += x;
    sumsquares += x*x;
    printf("%f\n", x);
    if ( i+1 < n )
    {
        sum += y;
        sumsquares += y*y;
        printf("%f\n", y);
    }
}
double sample_mean = sum/n;
double variance = (sumsquares - n * sample_mean * sample_mean)/(n-1);
double sample_standard_dev = sqrt (variance);
printf("n %d sample mean %f sample standard deviation %f\n",
       n, sample_mean, sample_standard_dev);

}

gcc -lm normal.c
a.out 4
prompt% a.out 4
-1.250447
-0.114427
-1.878096
-1.941784
n 4 sample mean -1.296189 sample standard deviation 0.847361
% a.out 4 (numbers should be N(0,1)). Not very close.
-0.713306
-1.578731
1.323076
0.233974
n 4 sample mean -0.183747 sample standard deviation 1.247854
Closer.

```

How does this work? It's best explained by describing a related method. Here is the Box-Muller version of `two_normally`.

```
// Box-Muller method
// U,V random (0,1)
// X = sqrt(-2 ln U) cos(2 pi V), Y = sqrt(-2 ln U) sin (2 pi V)

void two_normally( double * x, double * y )
{
    double u, v, mul;

    u = drand48();
    v = drand48();

    mul = sqrt ( - 2 * log(u) );
    *x = mul * cos ( 2 * M_PI * v );
    *y = mul * sin ( 2 * M_PI * v );
}
```

This works as follows. The multiplier M is a function of U , taking values in $(0, \infty)$, and (as one can check) with probability density function

$$\frac{1}{2}e^{-m/2}$$

But this PDF is χ_2^2 . Also, $m = x^2 + y^2$. so $X^2 + Y^2$ has the PDF of χ_2^2 . Also, X, Y are independent. It follows that X, Y have the joint PDF of two independent $N(0, 1)$ variables, which is what was wanted.

In the Marsaglia version, $S = U^2 + V^2$ is (since if $|S| \geq 1$ S is discarded) uniformly distributed between 0 and 1. If we write $u = S \cos \theta$ and $v = S \sin \theta$ then

$$X = U \sqrt{-\frac{2 \ln S}{S}} = \sqrt{-2 \ln S} \cos \theta$$

and similarly Y , so the formula is a variation of the Box-Muller formula.